

Learning Reversible Stochastic Tensor Transformations: A Path Toward Quantum-Proof MultiModal Encryption

Samuel Cavazos
Chief Data Scientist, Alshival.AI

October 2026

Abstract

Protecting personal communication and sensitive medical and scientific data is a central motivation for encryption research. A tensor-algebraic perspective offers a common mathematical language for studying transformations of images, audio, and other structured data. Crucially, the source and transmitted carrier need not share a modality: text could be carried as an image, or an image as audio, while a corresponding inverse recovers the original data. We explore this possibility as a foundation for future multimodal privacy systems, beginning with reliable recovery in an image-to-image experiment. Our initial study learned forward and reconstruction maps independently, without guaranteed inversion or held-out evaluation. Here we introduce a structural constraint: reversible coupling blocks provide an explicit inverse, while training learns to approximate a prescribed stochastic endpoint. A compact, 44,296-parameter network targets a 32-step Gaussian process, transforms all RGBA channels, and uses unclipped floating-point payloads to avoid destructive byte conversion. Three runs of 2,000 optimizer updates use 647 training, 81 validation, and 81 test Pokémon sprites; the test set was previously examined during exploratory work. Mean test forward- and reverse-target absolute errors are 0.00924 and 0.03694, respectively; mean learned round-trip error is 1.46×10^{-7} . The minimum exact byte-recovery rate across tested runs and domains is 100.0% after rounding. Target approximation is learned, whereas round-trip recovery relies on constructed inversion and numerical precision. Experiments are limited to images. These results establish numerical recovery under the evaluated conditions. A separate proof sketch shows how a reversible learned carrier can preserve one-time-pad secrecy against classical and quantum observation. This provides a conditional route toward the future-system conjecture; confidentiality of the experimental Gaussian transform remains unestablished.

1 Theory: tensor maps, stochastic endpoints, and recovery

The motivating question is whether we can learn to transform data into another representation while retaining a reliable way to recover its original values. That representation could have a different apparent modality: a message could travel as an image, or an image as an audio signal. This separates what the data originally means from the form that carries it.

The overarching idea begins with an invertible map $F : \mathcal{X} \rightarrow \mathcal{Y}$ between tensor representations. A machine-learning model σ learns the forward map, while a model τ learns the reverse map:

$$\sigma(x) \approx F(x), \quad \tau(y) \approx F^{-1}(y), \quad \tau(\sigma(x)) \approx x.$$

The first two relations express agreement with the desired transformations; the third asks whether the learned models work together to recover the input. Two independently learned approximations need not be exact inverses of one another. Our initial Encrypter–Decrypter study [1] learned the directions independently. Here we constrain the models to share an explicit inverse: writing $\sigma = E_\theta$ and $\tau = D_\theta$, training adjusts both directions through shared parameters, while the architecture guarantees $D_\theta = E_\theta^{-1}$ in exact arithmetic.

For our image-to-image experiment, we choose F from a stochastic process that progressively attenuates an image and adds Gaussian noise. Once the realized noise field n is fixed, the 32-step endpoint is an invertible map $F_n(x) = 0.25x + n$, with inverse $F_n^{-1}(y) = (y - n)/0.25$. Supplying the same n in both directions gives a controlled target whose forward and inverse operations are known, so we can measure target approximation separately from recovery of the network’s own output.

Figure 1 shows these two learned directions using Mew: σ transforms the sprite and τ recovers its original pixel values.




Figure 1: Mew through the learned forward map σ and its constructed inverse τ , using the frozen seed-17 model trained for 2,000 updates on the 32-step target. Both directions receive the same saved noise field. The middle panel shows a clipped RGB preview; recovery uses the full floating-point RGBA payload. Every original RGBA byte is recovered after rounding. Mew is a training image chosen to illustrate the maps; Section 3 reports the held-out results.

This learning objective can be posed for other invertible maps, including maps between modalities; the stochastic process is our chosen test case. Its extension to a particular map requires suitable tensor encodings, model capacity, training data, and storage precision. Different target maps may also require different reversible architectures. The theory below develops this tensor-map viewpoint through composition, smooth changes of coordinates, and a polynomial specialization. Section 2 relates recovery to the longer-term encryption objective, and Section 3 reports the experiments. Multimodal operation and cryptographic security remain research goals; the present evidence concerns image transformations and numerical recovery.

1.1 A modality-independent formulation

A digital image is a table of numbers. Our Mew sprite has four values at each of 120×120 positions: red, green, blue, and transparency. Arranging these values into a tensor gives us a precise way to describe how a transformation changes them. Audio can likewise be indexed by time and channel, and video by frame, channel, and spatial position. Tensor representations let us reason about maps between these forms as well as maps within one form. Each modality still needs its own suitable encoding and model.

Let $\mathcal{X} \subseteq \mathbb{R}^{d_1 \times \dots \times d_k}$ be a set of data tensors. For digital data, $\mathcal{X}$ may be a finite lattice; it need not itself be a vector space. An image is one instance, with spatial axes and color channels. Waveforms, spectrograms, video, and token sequences admit other tensor encodings. This common representation permits a shared mathematical formulation, while architecture, shape handling, and output constraints remain modality dependent. Exact recovery of the source data additionally requires the initial tensor encoding to be lossless: a lossy embedding cannot be inverted merely because a later tensor transform is invertible. The same trained image network is not claimed to operate on arbitrary modalities without adaptation.

The original paper’s tensor-map language [1] becomes more precise when we distinguish the stored data from its ambient space. For our images,

$$V = \mathbb{R}^4 \otimes \mathbb{R}^H \otimes \mathbb{R}^W \cong \mathbb{R}^N, \quad N = 4HW = 57,600.$$

The tensor-product notation $\otimes$ records the channel, height, and width axes. The symbol $\cong$ expresses an equivalence of vector spaces: we can list the same coordinates in one long vector and reshape them back without losing information. An entire sprite is one point in this space, with a coordinate for each channel value at each pixel. Byte-valued sprites occupy a finite grid inside V ; differential geometry describes the smooth maps we extend to the surrounding real coordinates. Here the tensor-algebraic viewpoint concerns these tensor spaces and the composition of maps between them. Our learned maps can be nonlinear and need not preserve vector addition or tensor products.

The source and carrier can have different forms. Consider the text MEW. A proposed encoder could represent its exact text bytes through the pixel values of an image; the corresponding decoder would recover MEW

from those values. The image need not display the letters or depict the message’s meaning. Similarly, an encoder could carry Mew’s RGBA pixel values in audio samples, with a decoder reconstructing the original sprite from the waveform. The sound need not describe Mew. The aim is to preserve the source information while changing its outward form, as illustrated in Figure 2.




Figure 2: Proposed transformations between modalities. The middle object carries the source information in a different form; the decoder recovers the original data, rather than generating a related description or approximation. Any noise, format metadata, and future secret-key inputs required by the design must also be available to the decoder. These are conceptual examples, not outputs from the current image-to-image model.

For source modality A and carrier modality B , this extends the recovery objective to

$$D_{B \rightarrow A}(E_{A \rightarrow B}(x, n), n) = x, \quad x \in \mathcal{X}_A, \quad (1)$$

where n denotes any conditioning noise shared by the two directions. For fixed n , distinct source inputs must remain distinguishable in the stored carrier. Its capacity and precision must therefore accommodate the admitted source data and the agreed format metadata. Lossless source encoding, compatible tensor shapes, and storage that retains the encoded distinctions are necessary; a modality change cannot restore information lost through lossy compression. Equation 1 is a design objective for these extensions.

Changing the carrier format alone does not require machine learning. The learning opportunity is to shape the carrier’s visual or acoustic properties while retaining recoverability. For example, the target might be an ordinary-looking image rather than an arbitrary array of pixels. Related neural data-hiding work, such as HiDDeN, learns to embed messages in cover images and recover them with a decoder [4]. Whether a carrier resembles ordinary media, hides the source modality, or protects the content from an adversary requires separate evaluation. Our present model and experiments establish image-to-image numerical recovery; they do not yet implement these cross-modal designs.

Before introducing the maps, consider the two paths in Figure 3. A prescribed formula generates the training target; we call this formula the *teacher*. A neural network learns to approximate that target. Its decoder is constructed to undo the network’s own output, and we also test how it handles the teacher’s output. The same realized noise field is supplied wherever n appears, including during decoding. It is saved receiver input, not something the decoder must guess.

Write a stochastic family as $T_\omega : \mathcal{X} \rightarrow \mathcal{Y}$, where ω specifies a noise realization. For fixed ω , it is an ordinary deterministic map. Our learned encoder and decoder satisfy the design objectives

$$E_\theta(x, n) \approx T_\omega(x), \quad D_\theta(E_\theta(x, n), n) \approx x, \quad (2)$$

where n summarizes the realized noise. The decoder is the encoder’s constructed inverse, with shared parameters. Section 1.4 develops the geometric and algebraic meaning of this inverse. Because Gaussian outputs are unbounded, $\mathcal{Y}$ contains real-valued tensors and is larger than the bounded image domain.

Same original x and saved noise n in both paths




Figure 3: Two paths through the experiment. The upper path creates the prescribed target and tests its reconstruction; the lower path creates, stores, and reverses the learned output. Forward error compares the paths’ outputs, reverse error tests the decoder on the teacher’s output, and round-trip error tests recovery from the network’s own output. The teacher’s 32 stochastic steps are represented by their endpoint formula; the learned encoder contains four coupling blocks.

1.2 The 32-step stochastic process

The original process corrupted RGB while preserving alpha, leaving the sprite’s silhouette directly available [1]. We now transform all four RGBA channels so that recovery includes transparency and no unchanged alpha mask is carried through the payload. This removes that particular source of leakage; it does not establish confidentiality. Each step slightly shrinks the current values and adds a fresh field of random perturbations. For normalized input $x_0 = x \in [0, 1]^{4 \times H \times W}$, this is

$$x_t = \sqrt{1 - \beta} x_{t-1} + \sqrt{\beta} \varepsilon_t, \quad \varepsilon_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, I). \quad (3)$$

Here iid means independent and identically distributed: each step draws a fresh noise field from the same distribution. The notation $\mathcal{N}(0, I)$ specifies Gaussian noise with mean zero and identity covariance, so its entries are independent and each has unit variance. The parameter β controls how much the current values shrink and how strongly the new noise enters.

The original paper used 1,000 steps with a varying noise schedule, leaving a signal coefficient of approximately 5.07×10^{-12} before clipping and byte quantization [1]. Such extreme attenuation made numerical recovery fragile even before those lossy storage operations. Our subsequent exploratory study used a constant schedule with a 16-step signal coefficient of 0.5. The present experiment retains that exploratory per-step strength $\beta = 1 - 0.5^{2/16} \approx 0.082996$ and increases T from 16 to 32. Set $a_T = (1 - \beta)^{T/2}$. Then $a_{16} = 0.5$ and $a_{32} = 0.25$. The integer T counts stochastic-process steps, independently of training updates and neural-network blocks. Increasing it from 16 to 32 reduces the residual signal under this fixed schedule while retaining a practical numerical inverse. Step counts alone do not compare corruption strength across different schedules.

Figure 4 makes this recurrence concrete using Mew. Each panel continues the same noise trajectory: we scale the current tensor and add the next independent Gaussian increment. The coefficient on the original image decreases from about 0.958 after one step to 0.25 after 32 steps. This is the teacher’s forward process, illustrated independently of the learned network.

Every earlier noise contribution is also scaled by the steps that follow it. Their weighted sum is the accumulated noise n_T . Independent Gaussian contributions combine into another Gaussian, giving a direct description of the endpoint distribution.

Proposition 1 (Gaussian endpoint). *In exact real arithmetic,*

$$x_T \mid x_0 \sim \mathcal{N}(a_T x_0, (1 - a_T^2)I), \quad x_T = a_T x_0 + n_T. \quad (4)$$

Proof. Unrolling the recurrence gives $n_T = \sum_{j=1}^T \sqrt{\beta}(1 - \beta)^{(T-j)/2} \varepsilon_j$. Independent Gaussian increments yield zero mean and covariance $\beta \sum_{j=0}^{T-1} (1 - \beta)^j I = [1 - (1 - \beta)^T]I$. $\square$




Figure 4: Mew at steps 0, 1, 2, 5, 16, and 32 along one reproducible stochastic trajectory, with $\beta \approx 0.082996$. Labels show the remaining signal coefficient, not a percentage of recoverable information. The original is shown with transparency on a checkerboard; later panels show opaque RGB previews clipped only for display. All four RGBA channels evolve in the underlying unclipped floating-point state. Mew is a training-set image chosen to explain the process; this figure is separate from the held-out reconstruction examples.

Thus the 32-step teacher is $y = 0.25x + n$, with $n \sim \mathcal{N}(0, 0.9375I)$. We sample this endpoint directly, avoiding the cost of simulating all 32 increments. This matches the endpoint distribution, not an individual unsaved noise trajectory. Given its realized n , the teacher has the analytic inverse $x = (y - n)/0.25$ and needs no machine learning. We deliberately choose a process with a known answer so we can measure what the network learns. This is a controlled test of target approximation within a reversible architecture, not evidence that a neural network is necessary or superior for this Gaussian transformation. Diffusion-inspired endpoint learning is distinct from sampling a plausible image from a learned distribution [5, 6].

1.3 Constructed inversion and numerical transport

Our initial paper took a *purely learned approach to forward and reverse mapping*: two separately trained U-Net models learned the prescribed corruption and reconstruction from paired examples [1]. “Purely learned” describes those endpoint maps; the stochastic process providing their targets was specified in advance. Reconstruction on the training collection showed that the models could fit those examples, but left recovery of unseen images unresolved. The reported evaluation did not establish either success or failure on a held-out collection.

Making an image look noisy is easy; retaining a reliable way back to its exact pixel values is the harder design requirement. Independently learning the two directions does not guarantee that one undoes the other. An unconstrained forward network can also discard distinctions between inputs that no decoder can restore. We therefore introduce a structural constraint: each transformation block comes with a known undo operation. Training determines how closely the resulting map matches the stochastic target, while the architecture provides its inverse in exact arithmetic. This is a change in method, not simply a longer training run.

The obstacle is easiest to see when a network computes an adjustment from the original image and then

changes every channel. To subtract that adjustment later, we may need the original image—the very thing we are trying to recover. Our channel split breaks this circular dependency: hold one group unchanged, use it and the saved noise to calculate an adjustment to the other group, and then switch roles. During reversal, the unchanged group is still available, so we can calculate exactly the same adjustment again and subtract it.

For Mew, the first group contains red and green, and the second contains blue and alpha. Both cover the entire image. The first block keeps red and green available while changing blue and alpha; the next keeps the resulting blue and alpha values available while changing red and green. Working backward restores each group’s earlier values when they are needed. Figure 5 shows this sequence. The split is one convenient way to construct a reversible learned map; the Gaussian teacher in Section 1.2 already has an analytic inverse without a split. The coupling construction cannot repair information lost later through clipping or coarse quantization, which we address below.

Formally, write the group held unchanged in a block as u and the group being updated as v . Each continuous coupling block applies

$$u' = u, \quad v' = sv + t_\theta(u, n), \quad v = \frac{v' - t_\theta(u', n)}{s}, \quad s \neq 0. \quad (5)$$

Alternating which half is updated and composing such blocks gives an explicit inverse by reversing their order, following the coupling principle used in Real NVP [2]. This algebraic property holds before training. Learning controls target-map fidelity, not the existence of the inverse.

Because $u' = u$, the decoder can run the same conditioner network on (u', n) to reproduce the forward adjustment, subtract it, and divide by s . The conditioner itself need not be inverted. In our model, $s = 0.5$ and four blocks alternate the groups’ roles, so both groups are changed twice. Decoding reverses the block order, restoring the values needed for each preceding operation. These four network blocks approximate the endpoint of the 32 stochastic steps in Section 1.2; the two counts refer to different operations.

An explicit inverse is useful only if storage retains the values it needs. The original pipeline clipped its noisy RGB endpoint to $[0, 1]$ and rounded it to 8-bit PNG values [1]; distinct floating-point values could therefore become the same stored value. We instead serialize the Gaussian payload as float32 NumPy arrays without clipping and decode from those arrays and the saved conditioning field. This avoids the earlier destructive conversion, although float32 still has finite precision. PNGs displayed as encrypted images are clipped RGB previews only, not decodable substitutes for the stored RGBA payload. For example, clipping sends both 1.2 and 1.8 to 1.0. Once only 1.0 is stored, the decoder cannot tell which value preceded it, even with a perfectly invertible network. This loss comes from storage, not from a failure to train the decoder long enough.

Finite arithmetic introduces reconstruction error. Normalizing byte values divides them by 255, so adjacent levels are $1/255$ apart. For original byte-valued inputs, error below half that spacing ($1/510$), allowing for input rounding, suffices to recover the byte grid by nearest rounding. We measure this directly rather than assuming it. The teacher inverse amplifies endpoint errors by $1/a_T = 4$; for example, a teacher-output error of 0.01 becomes a recovered input error of 0.04. Small errors can nevertheless leave the rounded byte unchanged: a recovered value of 127.99999 in byte units rounds back to 128. The learned composition also depends on its conditioner sensitivities. Exact-arithmetic invertibility therefore does not imply universal bitwise portability across devices or library versions.

These changes address distinct limitations: the schedule controls attenuation, coupling supplies an inverse, and the payload format preserves the numerical values needed to use it. Their combined recovery results cannot be attributed to channel splitting alone. Evaluation on images excluded from fitting, described in Section 3, separately examines target approximation and round-trip recovery beyond the training collection.

A. Split the channels; each group covers the whole image



B. One block: keep u available to reproduce the adjustment



C. Alternate groups across the four network blocks




Figure 5: Constructed inversion using Mew’s channel groups. (A) Original channel values are displayed in grayscale, including alpha; both groups cover every pixel. (B) One block holds u unchanged while adjusting v . The decoder uses the same conditioner and noise to reproduce the adjustment. The scalar values are an illustrative arithmetic example, not measured network predictions. (C) Alternating the groups changes all four channels; reversing the block order undoes the composition, up to floating-point error. A group is held unchanged only within its current block.

1.4 Geometry and algebra of reversible maps

Imagine nearby data points drawn on a coordinate grid. A reversible map can bend the grid and change distances while keeping every point distinguishable. Figure 6 illustrates this with two coordinates:

$$C(u, v) = (u, 0.5v + 0.3u^2), \quad C^{-1}(u', v') = (u', 2(v' - 0.3u'^2)).$$

The map compresses the vertical direction and shifts each vertical line by an amount determined by its horizontal coordinate. The inverse straightens the grid. For an image, the coordinates represent channel values across the whole tensor; this picture illustrates a deformation of data space.

Differential geometry: smooth changes of coordinates. Fix the noise field n . The implemented conditioners combine affine convolutions with the smooth activation $\text{SiLU}(z) = z/(1 + e^{-z})$. Their coupling maps therefore have smooth forward and inverse operations on V in exact real arithmetic. Such a map is called a *diffeomorphism*. The coupling construction makes this property explicit [2]. Zoom in on a tiny patch of the grid in Figure 6: the Jacobian describes its local stretching and shearing. The magnitude of its determinant gives the local area scale,




Figure 6: A two-coordinate coupling map and its inverse. The same two points, A and B, are tracked through all three panels. Horizontal grid lines become curves, vertical separation is halved, and every point remains recoverable. Both directions are polynomial maps in this example.

or volume scale in higher dimensions. For a block updating $q = N/2$ coordinates, its Jacobian is

$$J_C(u, v) = \begin{pmatrix} I_q & 0 \\ \partial_u t_\theta(u, n) & sI_q \end{pmatrix}, \quad \det J_C = s^q \neq 0. \quad (6)$$

In the two-coordinate example, the determinant is 0.5: area is halved while distinct points remain distinguishable. A nonzero determinant shows that no local direction is collapsed. Global invertibility follows from the explicit inverse in Equation 5; the determinant alone would establish only local invertibility.

With four blocks and $s = 0.5$, the full encoder satisfies $|\det J_{E_{\theta,n}}| = 0.5^{4q} = 0.25^N$, where $E_{\theta,n}(x) = E_\theta(x, n)$. Thus the model contracts local volume while preserving an inverse. The determinant’s magnitude gives the product of local stretching factors; individual directions can still be sensitive to perturbations. This explains why numerical error requires measurement even when a smooth inverse exists.

Composition and the algebraic viewpoint. Composition means applying one map after another. The notation is read from right to left: $C_2 \circ C_1$ applies C_1 first, then C_2 . Undoing that sequence requires undoing C_2 first, then C_1 . Writing the four blocks as $C_1, \dots, C_4$, with fixed n , gives

$$E_{\theta,n} = C_4 \circ C_3 \circ C_2 \circ C_1, \quad D_{\theta,n} = C_1^{-1} \circ C_2^{-1} \circ C_3^{-1} \circ C_4^{-1}, \quad (7)$$

and $D_{\theta,n} \circ E_{\theta,n} = E_{\theta,n} \circ D_{\theta,n} = \text{id}_V$ in exact arithmetic, where the identity map id_V leaves every input unchanged. An endomorphism maps a space into itself while respecting the structure under discussion; an automorphism additionally has an inverse respecting that same structure. Here the structure is smoothness: in the category of smooth manifolds, the encoder is an automorphism of the ambient tensor space V . This gives the original endomorphism viewpoint a specific setting. The bounded set of source images is only a subset of that space.

Algebraic geometry provides a related setting when the coordinate maps are polynomials [3]. If the conditioner $t_\theta(u, n)$ is polynomial in u for fixed n , then the coupling and its explicit inverse are polynomial, so the block is a polynomial automorphism of real affine space. The two-coordinate example above has this form, as does the fixed-noise Gaussian teacher $x \mapsto 0.25x + n$. The implemented SiLU conditioners supply smooth maps rather than a polynomial parameterization. Polynomial conditioners offer an algebraic variant for future experiments.

These distinctions also guide modality changes. Full continuous tensor spaces related by a diffeomorphism must have the same real dimension. A larger carrier space can instead contain an embedded copy of the source, with recovery defined on that copy. For finite byte encodings, the corresponding requirement is enough distinct stored carrier states for all admitted inputs. In each setting, the geometric or algebraic inverse addresses recovery; Section 2 supplies the separate secrecy argument for a keyed composition.

2 Connection to the pursuit of post-quantum encryption

2.1 What quantum resistance would require

Recovery and confidentiality answer different questions. Recovery asks whether the intended receiver can get the original data back. Confidentiality asks whether someone without the required secret can learn about it. In this prototype, a second person given the model, payload, and saved noise can run exactly the same decoder as the intended receiver. Accurate recovery therefore does not establish a distinction between authorized and unauthorized access.

A useful way to think about confidentiality is a guessing experiment. An observer knows two possible messages of equal length and receives a protected version of one, chosen with equal probability. Without a revealing observation, the observer can only guess with success probability $1/2$; indistinguishability asks whether seeing the protected version improves that guess. Perfect secrecy allows no improvement, while computational security requires any advantage available within the specified resource limits to be negligible as the security parameter grows.

The long-term objective is a system with efficient authorized inversion and computationally difficult unauthorized recovery. Quantum resistance concerns the capabilities of the adversary, not the absence of an inverse. Shor’s algorithms address factoring and discrete logarithms [7]; Grover’s algorithm supplies a quadratic speedup for unstructured search [8]. Neither says that every invertible map is easily attacked, nor that visually noisy or unfamiliar maps resist quantum computation. NIST’s post-quantum program evaluates explicit cryptographic constructions and security assumptions against classical and quantum attacks [9].

Our prototype is efficiently inverted by a classical computer given its public weights and noise field. The reproducible random generator is not a cryptographic key generator; its seed is experiment metadata. We release conditioning fields to reproduce recovery and make no secrecy claim for them. Even a hidden Gaussian field would not by itself establish perfect secrecy: the teacher’s conditional ciphertext distributions have input-dependent means. Residual structure and distributional leakage require analysis beyond how noisy a rendered image looks.

2.2 A future-system conjecture and a proof sketch

The longer-term aim is to create that distinction through a specified secret-key mechanism and to evaluate it against both classical and quantum adversaries. Changing a message’s outward form, as in the text-to-image example, adds a possible concealment objective: an observer might not recognize the source modality. Demonstrating that a carrier is indistinguishable from ordinary media is a steganographic question, distinct from whether its hidden content is cryptographically confidential. Neither property follows from representing the data as a tensor or from changing its file format. The following conjecture guides the search for a keyed construction. One route is to combine an established secrecy mechanism with a learned carrier; the sketch below gives a conditional result for this approach.

Conjecture 1 (Research hypothesis; unproven). *Under suitable computational hardness assumptions, a future publicly specified key-conditioned family of learned, information-preserving tensor transformations can support efficient authorized recovery and confidentiality against polynomial-time classical and quantum adversaries.*

For a computationally secure implementation, a proposed construction must specify a security parameter, secret-key distribution, public randomness or nonce rules, ciphertext format, and an indistinguishability experiment with precisely defined oracle access. At minimum, one should consider quantum computation with classical chosen-plaintext queries; stronger quantum-query models require separate claims. Models and architecture should be assumed public. Authentication, key exchange, chosen-ciphertext resistance, and side-channel behavior are additional design obligations, none supplied by reconstruction accuracy.

Sketch of Proof: preserving secrecy through a learned carrier. First protect the message, then change the form that carries it. If the protected bits reveal no information about the message, expressing those same bits as an image or waveform cannot create information about the message. The receiver reverses the carrier transformation and then removes the protection using a secret key.

Fix a public message length L and a lossless representation $M \in \{0,1\}^L$. The sender and receiver share a uniformly random secret pad $K \in \{0,1\}^L$, independent of the message and the adversary’s prior information. Every transmission uses a fresh pad. Let F_θ be a public, efficiently computable learned map from L -bit strings to serialized image or audio carriers. Its parameters are fixed independently of the secret message and pad. Assume

an efficient inverse G_θ satisfying $G_\theta(F_\theta(c)) = c$ for every $c \in \{0, 1\}^L$, including the effects of serialization. The transmitted carrier is

$$C = M \oplus K, \quad Y = F_\theta(C), \quad M = G_\theta(Y) \oplus K, \quad (8)$$

where $\oplus$ is bitwise exclusive OR. The final identity establishes correctness. The carrier can have any modality with sufficient capacity for these ciphertexts.

For example, a four-bit message 1010 and pad 0110 give $1010 \oplus 0110 = 1100$. After reversing the carrier map, applying the same pad recovers the message: $1100 \oplus 0110 = 1010$. XOR flips a message bit wherever the pad has a 1; applying those same flips twice restores the original bits.

For any fixed message, each possible ciphertext corresponds to exactly one pad, and all pads are equally likely. The one-time-pad secrecy argument [10] therefore gives, for every m, c ,

$$\Pr[C = c \mid M = m] = \Pr[K = m \oplus c] = 2^{-L}. \quad (9)$$

Thus every message of length L induces the same distribution of C . Applying the fixed carrier map preserves that equality:

$$\Pr[Y = y \mid M = m] = \sum_{c: F_\theta(c)=y} 2^{-L}, \quad (10)$$

which is independent of m . An observer may know and reverse F_θ , but this yields only the uniformly masked ciphertext C .

The equality also covers quantum computation on the observed classical carrier. We write its distribution as a diagonal density operator,

$$\rho_m = \sum_y \Pr[Y = y \mid M = m] |y\rangle\langle y|, \quad (11)$$

where $|y\rangle\langle y|$ is the outer product of the basis vector $|y\rangle$ representing output y with its conjugate transpose $\langle y|$. This outer product marks the diagonal position corresponding to output y : it has a 1 in that position and zeros elsewhere. Weighting each outer product by its probability therefore places the output probabilities along the diagonal of ρ_m . Since those probabilities are the same for every message, we have $\rho_m = \rho_{m'}$ for every pair of messages. Identical input states produce identical outcome distributions under every quantum computation and measurement. The carrier therefore adds no distinguishing advantage, even with unlimited computation. This argument concerns observation of classical transmissions with independent pads; it assumes no access to the secret pad or a decryption oracle. Public metadata must be fixed or generated from the ciphertext and independent public randomness, so that the whole transmitted record satisfies the same independence condition. Message length is public.

This construction establishes a conditional secrecy-preserving route toward the conjecture. The pad supplies perfect secrecy, and the learned map supplies the carrier representation. Its cost is an independently shared secret pad as long as each message. Replacing the pad with a short pseudorandom seed would require a separate computational security argument. Authentication and resemblance to ordinary images or audio also require separate mechanisms or evidence. The Gaussian noise used in our experiments is not a one-time pad, and measured round-trip accuracy does not establish the assumed exact carrier inverse for every possible bit string. Constructing such a carrier, or bounding its decoding error, is the next implementation step.

3 Architecture, training, and Pokémon experiments

3.1 Model and training protocol

Replacing the original pair of independently trained U-Nets [1], the network has four alternating channel-coupling blocks with shared forward and inverse parameters. Each conditioner uses a 3×3 convolution from six channels (two unchanged image channels plus four noise channels) to 32 channels, SiLU, a second 3×3 32-channel convolution, SiLU, and a 1×1 output convolution to two channels. There are 44,296 trainable parameters. Each half is updated twice; the fixed scale $s = 0.5$ makes its total multiplicative attenuation 0.25. This builds the teacher’s known overall scaling into the architecture: training learns the additive adjustments needed to match its endpoint. Supplying this part of the target structure in advance is the model’s known-scale inductive bias.

The original models classified each RGB value into one of 256 byte bins [1]. Because our transport retains continuous-valued RGBA tensors, we use regression against the floating-point targets. The loss is the sum of

forward and reverse mean squared errors:

$$\mathcal{L} = \mathbb{E} [\|E_{\theta}(x, n) - T_{\omega}(x)\|_{\text{mean}}^2 + \|D_{\theta}(T_{\omega}(x), n) - x\|_{\text{mean}}^2]. \quad (12)$$

Three independent runs use seeds 17, 29, and 43, each with 2,000 optimizer updates, effective batch size 16, and microbatch size four. AdamW uses learning rate 3×10^{-4} , weight decay 10^{-5} , and gradient clipping at one. These are optimizer updates, not 2,000 passes through the dataset. All coupling arithmetic uses FP32 with TF32 disabled and deterministic cuDNN settings.

The original study fitted and evaluated the same paired collection [1]. To distinguish fitting that collection from performance on images excluded from training, we use separate training, validation, and test sets. Training fits parameters, validation selects checkpoints, and test evaluation measures the selected models. The 809 native 120×120 sprites come from the *Pokemon sprite images* dataset shared by Kaggle user yehongjiang [11] and retain their RGBA channels. The inherited split contains 647 training, 81 validation, and 81 test images, grouping decoded duplicates, name variants, and available evolution links. The supplied evolution metadata are incomplete. Training samples sprites with replacement and fresh noise. Validation every 100 updates selects checkpoints by round-trip MAE above a 10^{-6} tolerance, followed by summed forward/reverse target MAE. All three checkpoints are frozen and hashed before this campaign’s test evaluation. The same test set was previously examined during exploratory work; consequently this is a follow-up evaluation, not a previously untouched confirmatory holdout. Further tuning would require fresh confirmatory data.

Uniform static was not included in training because the earlier fixed-budget comparison did not improve with it. That finding does not rule out structured synthetic data or exposure-matched augmentation. Independent test streams of 32 static images and nine structured probes supplement the 81 Pokémon test images.

3.2 Losses, target fidelity, and reconstruction

An accurate round trip can coexist with an imperfect match to the teacher. Consider a scalar example with $x = 0.8$ and saved noise $n = 0.1$. The teacher gives $T(x, n) = 0.25x + n = 0.30$. Suppose a learned map instead gives $E(x, n) = 0.25x + n + 0.01 = 0.31$, with constructed inverse $D(z, n) = (z - n - 0.01)/0.25$. Then

$$\begin{aligned} \text{forward error:} & \quad |0.31 - 0.30| = 0.01, \\ \text{reverse error:} & \quad |D(0.30, 0.1) - 0.8| = |0.76 - 0.8| = 0.04, \\ \text{round-trip error:} & \quad |D(0.31, 0.1) - 0.8| = 0. \end{aligned}$$

The map perfectly reverses its own output even though it differs from the teacher’s transformation.

Forward error measures agreement with the prescribed teacher; reverse error decodes teacher outputs. Round-trip error instead decodes the network’s own serialized output. Their separation prevents built-in invertibility from being misreported as learned target fidelity. MAE is normalized to original pixel units on $[0, 1]$, including all RGBA entries. Foreground RGB errors use the original nonzero-alpha mask and are included in the accompanying results.

Our saved untrained controls make the same distinction experimentally. Table 1 compares the randomly initialized networks with their selected trained checkpoints on the same validation collection. Every image was already recovered exactly after rounding before training, although the target errors were large. Training substantially reduced those target errors. The round-trip property comes from the architecture; training teaches its additive functions to match the prescribed process.

Model state	Forward MAE	Reverse MAE	Exact RGBA images (%)
Before training	0.77634	3.11343	100.0
After training	0.00924	0.03691	100.0

Table 1: Before and after training on the 81 validation Pokémon images. Entries are means across seeds 17, 29, and 43. “After” uses each seed’s validation-selected checkpoint. Exact RGBA images counts whole images with every channel byte recovered after rounding; all three seeds attain 100% both before and after training.




Figure 7: Training loss averaged over successive 100-update intervals, and validation target error. Each curve follows one training seed.

Seed	Best update	Forward MAE	Reverse MAE	Round-trip MAE	Exact RGBA images (%)
17	2000	0.00934	0.03736	1.5×10^{-7}	100.0
29	2000	0.00918	0.03666	1.43×10^{-7}	100.0
43	2000	0.00922	0.03679	1.45×10^{-7}	100.0

Table 2: Results on the 81 test Pokémon images, shown separately for each training seed. MAEs use normalized RGBA values. Exact RGBA images is the percentage of entire images recovered byte-for-byte after rounding.

Mean foreground-only RGB reverse-target MAE at 32 steps was 0.04199. The minimum exact RGBA image-recovery rate across all three seeds and all three test domains was 100.0%. The largest floating-point round-trip error across those runs was 2.92×10^{-6} . These measurements describe the evaluated images and runtime; they are not a universal finite-precision guarantee. Per-image and per-domain results, untrained validation controls, selected update numbers, environment records, source code, and frozen checkpoints are available in the reproducibility repository [12]. Seed variation describes training variation and is not a population confidence interval.

3.3 Comparison with the 16-step study

Moving to 32 steps produced a tradeoff: measured RGB correlation with the input was lower and forward-target agreement improved, while reverse-target agreement worsened. Exact byte recovery of the models’ own outputs remained 100% in the tested domains. These measurements describe different aspects of the method, so no single one establishes that every aspect improved.

The 16-step exploratory study followed the original paper and already used coupling and float transport. It used the same model capacity, training split, update budget, and seeds as the present experiment, with endpoint signal coefficient 0.5. The 32-step model is newly trained with coefficient 0.25. The following table reports mean and sample standard deviation across three seeds on the same test images. The comparison changes the teacher, so the metrics describe two task difficulties rather than a competition against an identical target. The mean

Steps	Signal	Forward MAE	Reverse MAE	Round-trip MAE
16	0.5	0.0164 ± 0.00012	0.0328 ± 0.00024	$7.58 \times 10^{-8} \pm 1 \times 10^{-9}$
32	0.25	$0.00924 \pm 8.6 \times 10^{-5}$	0.0369 ± 0.00037	$1.46 \times 10^{-7} \pm 3.5 \times 10^{-9}$

Table 3: Target and round-trip MAEs on the same 81 test Pokémon images. Entries are mean $\pm$ sample standard deviation across three training seeds. Signal is the teacher’s coefficient on the original input. The two schedules define different targets.

forward-plus-reverse target MAE decreased from 0.04924 at 16 steps to 0.04618 at 32 steps. Separately, forward-target MAE changed from 0.01641 to 0.00924, while reverse-target MAE changed from 0.03283 to 0.03694. The two directions should therefore also be assessed separately.




Figure 8: Post-hoc descriptive correlation between original RGB and each learned output, using full float payloads and clipped previews. Error bars show seed standard deviation. Lower correlation does not establish confidentiality or rule out nonlinear identification or reconstruction attacks.

Mean preview correlation is 0.1195 for 16 steps and 0.0515 for 32 steps. Noise fields use matching underlying standard-normal streams for the comparison. Both positive and negative findings are retained; no claim that all recognizable structure is eliminated follows from these summaries.

3.4 Actual image examples and resource costs

Figure 9 illustrates one predetermined test image; the appendix contains all first 25 in manifest order, without selection for appearance or reconstruction quality. Checkerboards display original and reconstructed transparency. The encrypted preview shows opaque clipped RGB; the saved float payload also contains transformed alpha.




Figure 9: Amoonguss, the first test image in manifest order (seed 17). Float round-trip MAE is 1.51×10^{-7} ; all RGBA bytes are recovered exactly after rounding. The center panel is a clipped preview; recovery uses the full float payload and saved noise. This example illustrates numerical recovery, not confidentiality.

On the RTX 5070 with 12 GB VRAM, training took 2.4–2.5 minutes per run, excluding validation and checkpoint overhead. Peak PyTorch tensor allocation was 380–380 MiB; driver memory and other processes are excluded. Timing describes this shared workstation. The float32 RGBA payload and noise field each occupy 230,400 raw bytes, excluding container overhead and shared model weights. Replacing the noise field with a seed would require generator and stream-position metadata and separate security analysis.

This small experiment demonstrates learned target approximation and reliable numerical round trips under the tested conditions. It establishes neither an advantage over the analytic inverse nor cross-device bitwise compatibility. The text-to-image and image-to-audio paths in Figure 2 are concrete next experiments: their first recovery criterion would be exact source bytes after serialization of the carrier. Carrier realism and resilience to format conversions would need additional tests, alongside a specified keyed construction and independent attacker evaluations for security claims.

A All 25 predetermined test examples

Every panel below uses the seed-17 model and the saved float payload. Encrypted panels are display-only RGB previews. All 25 examples were chosen by manifest order before examining this campaign's outputs.




Figure 10: Test examples 1–5.




Figure 11: Test examples 6–10.




Figure 12: Test examples 11–15.




Figure 13: Test examples 16–20.




Figure 14: Test examples 21–25.

References

- [1] S. Cavazos. Learned Encrypter-Decrypter Models for Diffusion-Like Image Encryption: A Machine-Learning Approach Based on Stochastic Processes. Alshival.AI, original manuscript and implementation. <https://github.com/Alshival-Ai/alshicrypt-multimodal>.
- [2] L. Dinh, J. Sohl-Dickstein, and S. Bengio. Density estimation using Real NVP. ICLR, 2017. <https://arxiv.org/abs/1605.08803>.
- [3] A. van den Essen. Polynomial Automorphisms and the Jacobian Conjecture. Séminaires et Congrès, 2:55–81, 1997. <https://smf.emath.fr/publications/polynomial-automorphisms-and-jacobian-conjecture>.
- [4] J. Zhu, R. Kaplan, J. Johnson, and L. Fei-Fei. HiDDeN: Hiding Data with Deep Networks. ECCV, 2018. <https://arxiv.org/abs/1807.09937>.
- [5] J. Ho, A. Jain, and P. Abbeel. Denoising Diffusion Probabilistic Models. NeurIPS, 2020. <https://arxiv.org/abs/2006.11239>.
- [6] Y. Song et al. Score-Based Generative Modeling through Stochastic Differential Equations. ICLR, 2021. <https://arxiv.org/abs/2011.13456>.
- [7] P. W. Shor. Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer. SIAM Journal on Computing, 1997. <https://arxiv.org/abs/quant-ph/9508027>.
- [8] L. K. Grover. A fast quantum mechanical algorithm for database search. STOC, 1996. <https://arxiv.org/abs/quant-ph/9605043>.
- [9] NIST. Post-Quantum Cryptography: Frequently Asked Questions. Accessed October 2026. <https://csrc.nist.gov/Projects/Post-Quantum-Cryptography/faqs>.
- [10] C. E. Shannon. Communication Theory of Secrecy Systems. Bell System Technical Journal, 28(4):656–715, 1949. <https://doi.org/10.1002/j.1538-7305.1949.tb00928.x>.
- [11] yehongjiang. Pokemon sprite images. Kaggle dataset. <https://www.kaggle.com/datasets/yehongjiang/pokemon-sprites-images>.
- [12] S. Cavazos. Alshicrypt multimodal, October 2026: source code, frozen checkpoints, and experimental records. <https://github.com/Alshival-Ai/alshicrypt-multimodal-oct-2026>.